

# PROBLEM STATEMENT

## Context:

Porter is India's Largest Marketplace for Intra-City Logistics. Leader in the country's \$40 billion intra-city logistics market, Porter strives to improve the lives of 1,50,000+ driver-partners by providing them with consistent earning & independence. Currently, the company has serviced 5+ million customers

Porter works with a wide range of restaurants for delivering their items directly to the people.

Porter has a number of delivery partners available for delivering the food, from various restaurants and wants to get an estimated delivery time that it can provide the customers on the basis of what they are ordering, from where and also the delivery partners.

This dataset has the required data to train a regression model that will do the delivery time estimation, based on all those features

### Data Dictionary

Each row in this file corresponds to one unique delivery. Each column corresponds to a feature as explained below.

market\_id : integer id for the market where the restaurant lies

created\_at : the timestamp at which the order was placed

actual\_delivery\_time : the timestamp when the order was delivered

store\_primary\_category : category for the restaurant

order\_protocol : integer code value for order protocol(how the order was placed ie: through porter, call to restaurant, pre booked, third part etc)

total\_items subtotal : final price of the order

num\_distinct\_items : the number of distinct items in the order

min\_item\_price : price of the cheapest item in the order

max\_item\_price : price of the costliest item in order

total\_onshift\_partners : number of delivery partners on duty at the time order was placed

total\_busy\_partners : number of delivery partners attending to other tasks

total\_outstanding\_orders : total number of orders to be fulfilled at the moment

## Importing libraries

```
In [1]: # Importing necessary libraries for data manipulation and analysis
import pandas as pd
import numpy as np
# Importing libraries for data visualization
import matplotlib.pyplot as plt
import seaborn as sns
# Importing libraries for handling datetime operations
from datetime import datetime
# Importing libraries for preprocessing and encoding
from sklearn.preprocessing import StandardScaler, OneHotEncoder
# Importing libraries for splitting data
from sklearn.model_selection import train_test_split
# Importing libraries for neural network
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
# Importing libraries for evaluation metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error
# Setting up visualization styles
sns.set(style='whitegrid')
plt.rcParams['figure.figsize'] = (10, 6)
# Ignoring warnings
import warnings
warnings.filterwarnings('ignore')
```



## Importing the data

```
In [2]: df = pd.read_csv(r"H:\Scaler\Deep learning\Porter NN Project\dataset.csv\dataset.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	market_id	created_at	actual_delivery_time	store_id	store_primary_category	order_protocol	total_items	su
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	df263d996281d984952c07998dc54358	american	1.0	4	
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	f0ade77b43923b38237db569b016ba25	mexican	2.0	1	
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	f0ade77b43923b38237db569b016ba25	NaN	1.0	1	
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	f0ade77b43923b38237db569b016ba25	NaN	1.0	6	
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	f0ade77b43923b38237db569b016ba25	NaN	1.0	3	

```
In [4]: df.shape
```

```
Out[4]: (197428, 14)
```

The dataset has 197428 rows and 14 columns

## Checking info of the features in the dataset

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            196441 non-null float64
1   created_at                           197428 non-null object
2   actual_delivery_time                 197421 non-null object
3   store_id                             197428 non-null object
4   store_primary_category               192668 non-null object
5   order_protocol                       196433 non-null float64
6   total_items                          197428 non-null int64
7   subtotal                             197428 non-null int64
8   num_distinct_items                   197428 non-null int64
9   min_item_price                       197428 non-null int64
10  max_item_price                       197428 non-null int64
11  total_onshift_partners               181166 non-null float64
12  total_busy_partners                  181166 non-null float64
13  total_outstanding_orders             181166 non-null float64
dtypes: float64(5), int64(5), object(4)
memory usage: 21.1+ MB
```

Checking for null values



```
In [6]: df.isnull().sum()
```

```
Out[6]: market_id          987
created_at          0
actual_delivery_time  7
store_id            0
store_primary_category 4760
order_protocol      995
total_items         0
subtotal            0
num_distinct_items  0
min_item_price      0
max_item_price      0
total_onshift_partners 16262
total_busy_partners  16262
total_outstanding_orders 16262
dtype: int64
```

```
In [7]: df.isnull().sum().sum()
```

```
Out[7]: 55535
```

There are a total of 55535 null values in the dataset

The market id column has 987 null values

The actual delivery time has 7 null values

The store\_primary\_category has 4760 null values

The order\_protocol has 995 null values

The total\_onshift\_partners , total\_busy\_partners , total\_outstanding\_orders has 16262 null values each respectively

```
In [8]: df.isna().sum()/df.shape[0]*100
```

```
Out[8]: market_id          0.499929
created_at          0.000000
actual_delivery_time  0.003546
store_id            0.000000
store_primary_category 2.411006
order_protocol      0.503981
total_items         0.000000
subtotal            0.000000
num_distinct_items  0.000000
min_item_price      0.000000
max_item_price      0.000000
total_onshift_partners 8.236927
total_busy_partners  8.236927
total_outstanding_orders 8.236927
dtype: float64
```

Converting data type of columns created\_at and actual\_delivery\_time to date time

```
In [9]: DFDT = ['created_at', 'actual_delivery_time']

for i in DFDT:
    df[i] = pd.to_datetime(df[i])
```



In [10]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                             196441 non-null  float64
1   created_at                             197428 non-null  datetime64[ns]
2   actual_delivery_time                   197421 non-null  datetime64[ns]
3   store_id                               197428 non-null  object
4   store_primary_category                 192668 non-null  object
5   order_protocol                         196433 non-null  float64
6   total_items                           197428 non-null  int64
7   subtotal                              197428 non-null  int64
8   num_distinct_items                    197428 non-null  int64
9   min_item_price                        197428 non-null  int64
10  max_item_price                         197428 non-null  int64
11  total_onshift_partners                 181166 non-null  float64
12  total_busy_partners                   181166 non-null  float64
13  total_outstanding_orders              181166 non-null  float64
dtypes: datetime64[ns](2), float64(5), int64(5), object(2)
memory usage: 21.1+ MB
```

Creating target column(Time taken)

```
In [11]: # Create a new column named 'time_taken' to store the difference in minutes
df['time_taken'] = (df['actual_delivery_time'] - df['created_at'])
```

In [12]: df.head()

Out[12]:

	market_id	created_at	actual_delivery_time	store_id	store_primary_category	order_protocol	total_items	su
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	df263d996281d984952c07998dc54358	american	1.0	4	
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	f0ade77b43923b38237db569b016ba25	mexican	2.0	1	
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	f0ade77b43923b38237db569b016ba25	NaN	1.0	1	
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	f0ade77b43923b38237db569b016ba25	NaN	1.0	6	
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	f0ade77b43923b38237db569b016ba25	NaN	1.0	3	

```
In [13]: # Extracting the total minutes from the 'time_taken' column
df['time_taken_minutes'] = df['time_taken'].dt.total_seconds() // 60
```

In [14]: df.head()

Out[14]:

	market_id	created_at	actual_delivery_time	store_id	store_primary_category	order_protocol	total_items	su
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	df263d996281d984952c07998dc54358	american	1.0	4	
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	f0ade77b43923b38237db569b016ba25	mexican	2.0	1	
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	f0ade77b43923b38237db569b016ba25	NaN	1.0	1	
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	f0ade77b43923b38237db569b016ba25	NaN	1.0	6	
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	f0ade77b43923b38237db569b016ba25	NaN	1.0	3	



## Feature Engineering and Data Preprocessing

```
In [15]: # Extracting hour and day of the week from 'created_at'
df['order_hour'] = df['created_at'].dt.hour
df['order_day_of_week'] = df['created_at'].dt.dayofweek # Monday=0, Sunday=6
```

```
In [16]: df.head()
```

```
Out[16]:
```

	market_id	created_at	actual_delivery_time	store_id	store_primary_category	order_protocol	total_items	subtotal
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	df263d996281d984952c07998dc54358	american	1.0	4	12.99
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	f0ade77b43923b38237db569b016ba25	mexican	2.0	1	12.99
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	f0ade77b43923b38237db569b016ba25	NaN	1.0	1	12.99
3	3.0	2015-02-03 21:21:45	2015-02-03 22:13:00	f0ade77b43923b38237db569b016ba25	NaN	1.0	6	12.99
4	3.0	2015-02-15 02:40:36	2015-02-15 03:20:26	f0ade77b43923b38237db569b016ba25	NaN	1.0	3	12.99

Dropping columns that aren't useful anymore

```
In [18]: df.drop(['time_taken', 'created_at', 'actual_delivery_time'], axis=1, inplace=True)
```

```
-----
KeyError                                Traceback (most recent call last)
Cell In[18], line 1
----> 1 df.drop(['time_taken', 'created_at', 'actual_delivery_time'], axis=1, inplace=True)
      3 df.info()

File D:\Users\india\anaconda3\lib\site-packages\pandas\util\_decorators.py:331, in deprecate_nonkeyword_arguments.<locals>.decorate.<locals>.wrapper(*args, **kwargs)
    325 if len(args) > num_allow_args:
    326     warnings.warn(
    327         msg.format(arguments=_format_argument_list(allow_args)),
    328         FutureWarning,
    329         stacklevel=find_stack_level(),
    330     )
--> 331 return func(*args, **kwargs)

File D:\Users\india\anaconda3\lib\site-packages\pandas\core\frame.py:5399, in DataFrame.drop(self, labels, axis, index, columns, level, inplace, errors)
    5251 @deprecate_nonkeyword_arguments(version=None, allowed_args=["self", "labels"])
    5252 def drop(
    5253     self,
    5254     labels=None,
    5255     axis=0,
    5256     index=None,
    5257     columns=None,
    5258     level=None,
    5259     inplace=False,
    5260     errors="raise",
    5261 ):
    5262     """
    5263     Drop labels from the index, columns, or row labels.
    5264     """
    5265     axis = self._get_axis_number(axis)
    5266     axis = self._convert_to_int(axis)
    5267     labels = _convert_to_list_like(labels)
    5268     labels = _convert_to_index_labels(labels, axis=axis)
    5269     if axis == 0 and index is None:
    5270         index = labels
    5271     if axis == 1 and columns is None:
    5272         columns = labels
    5273     if axis == 0:
    5274         new_index, indexer = self.index.delete(index=labels, inplace=False, errors=errors)
    5275         return self._reindex_with_indexers({"index": (new_index, indexer)},
    5276                                           allow_dups=True)
    5277     if axis == 1:
    5278         new_columns, indexer = self.columns.delete(labels, inplace=False, errors=errors)
    5279         return self._reindex_with_indexers({"columns": (new_columns, indexer)},
    5280                                           allow_dups=True)
    5281     if axis == 0 and columns is not None:
    5282         # Drop labels from the index, but not the columns
    5283         new_index, indexer = self.index.delete(index=labels, inplace=False, errors=errors)
    5284         return self._reindex_with_indexers({"index": (new_index, indexer)},
    5285                                           allow_dups=True)
    5286     if axis == 1 and index is not None:
    5287         # Drop labels from the columns, but not the index
    5288         new_columns, indexer = self.columns.delete(labels, inplace=False, errors=errors)
    5289         return self._reindex_with_indexers({"columns": (new_columns, indexer)},
    5290                                           allow_dups=True)
    5291     raise ValueError("No labels specified for dropping")
```

```
In [19]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                             196441 non-null  float64
1   store_id                              197428 non-null  object
2   store_primary_category                 192668 non-null  object
3   order_protocol                         196433 non-null  float64
4   total_items                           197428 non-null  int64
5   subtotal                              197428 non-null  int64
6   num_distinct_items                    197428 non-null  int64
7   min_item_price                        197428 non-null  int64
8   max_item_price                        197428 non-null  int64
9   total_onshift_partners                 181166 non-null  float64
10  total_busy_partners                    181166 non-null  float64
11  total_outstanding_orders               181166 non-null  float64
12  time_taken_minutes                     197421 non-null  float64
13  order_hour                            197428 non-null  int64
14  order_day_of_week                      197428 non-null  int64
dtypes: float64(6), int64(7), object(2)
memory usage: 22.6+ MB
```



## Handling Null values

In [20]: `df.isna().sum()`

```
Out[20]: market_id          987
store_id              0
store_primary_category 4760
order_protocol        995
total_items           0
subtotal              0
num_distinct_items    0
min_item_price        0
max_item_price        0
total_onshift_partners 16262
total_busy_partners    16262
total_outstanding_orders 16262
time_taken_minutes     7
order_hour             0
order_day_of_week      0
dtype: int64
```

```
In [22]: # Finding the number of unique values in each column
unique_values = {column: df[column].nunique() for column in df.columns}
# Displaying the unique values count for each column
for column, unique_count in unique_values.items():
    print(f"{column}: {unique_count}")
```

```
market_id: 6
store_id: 6743
store_primary_category: 74
order_protocol: 7
total_items: 57
subtotal: 8368
num_distinct_items: 20
min_item_price: 2312
max_item_price: 2652
total_onshift_partners: 172
total_busy_partners: 159
total_outstanding_orders: 281
time_taken_minutes: 274
order_hour: 19
order_day_of_week: 7
```

In [23]: `df1=df.dropna()`

In [24]: `df[df["store_id"]=="252a3dbaeb32e7690242ad3b556e626b"]`

Out[24]:

	market_id	store_id	store_primary_category	order_protocol	total_items	subtotal	num_distinct_items
52018	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	4	5950	3
52019	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	1	2735	1
52020	2.0	252a3dbaeb32e7690242ad3b556e626b	burger	3.0	2	2515	2
52021	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	2	3915	2
52022	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	1	2064	1
...	...	...	...	...	...	...	...
63432	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	1	1828	1
63433	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	3	4055	2
63434	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	1	1510	1
63435	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	1	1890	1
63436	6.0	252a3dbaeb32e7690242ad3b556e626b	american	5.0	2	2235	2

350 rows × 15 columns

## Checking whether mean or median is the right choice for Null

imputation



```
In [25]: df.groupby("market_id")["total_onshift_partners"].mean()
```

```
Out[25]: market_id
1.0      24.208854
2.0      62.590695
3.0      18.847580
4.0      60.464482
5.0      23.911045
6.0      44.929771
Name: total_onshift_partners, dtype: float64
```

```
In [26]: df.groupby("market_id")["total_onshift_partners"].median()
```

```
Out[26]: market_id
1.0      19.0
2.0      55.0
3.0      15.0
4.0      60.0
5.0      20.0
6.0      36.0
Name: total_onshift_partners, dtype: float64
```

```
In [27]: df.groupby("order_hour")["total_onshift_partners"].mean()
```

```
Out[27]: order_hour
0      27.933751
1      54.325601
2      67.995169
3      64.205588
4      44.996112
5      23.589613
6      13.421094
7      10.777778
8       0.000000
14     0.550000
15     2.141473
16     4.965949
17     7.757729
18    15.092275
19    32.199487
20    37.353387
21    30.325540
22    22.749043
23    20.274580
Name: total_onshift_partners, dtype: float64
```

```
In [28]: df.groupby("order_day_of_week")["total_onshift_partners"].mean()
```

```
Out[28]: order_day_of_week
0      42.084044
1      37.333062
2      40.067352
3      43.746503
4      48.602855
5      52.111917
6      45.943654
Name: total_onshift_partners, dtype: float64
```

```
In [29]: df.groupby(["market_id", "order_hour"])["total_onshift_partners"].mean()
```

```
Out[29]: market_id  order_hour
1.0              0      14.437811
              1      26.014145
              2      36.809734
              3      37.072227
              4      27.385254
              ...
6.0             19      30.744186
              20      40.627907
              21      31.200000
              22      23.806452
              23      18.000000
Name: total_onshift_partners, Length: 106, dtype: float64
```

## Mean Imputation



```
In [32]: # List of columns to impute
columns_to_impute = ['total_outstanding_orders', 'total_busy_partners', 'total_onshift_partners']
# Group by 'market_id' and 'order_hour'
grouped = df.groupby(['market_id', 'order_hour'])
# Impute missing values
for column in columns_to_impute:
    # Calculate the mean for each group and transform to align with the original Data
    df[column] = grouped[column].transform(lambda x: x.fillna(x.mean()))
```

```
In [33]: df
```

```
Out[33]:
```

	market_id	store_id	store_primary_category	order_protocol	total_items	subtotal	num_distinct_items
0	1.0	df263d996281d984952c07998dc54358	american	1.0	4	3441	4
1	2.0	f0ade77b43923b38237db569b016ba25	mexican	2.0	1	1900	1
2	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	1	1900	1
3	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	6	6900	5
4	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	3	3900	3
...	...	...	...	...	...	...	...
197423	1.0	a914ecef9c12fdb9bede64bb703d877	fast	4.0	3	1389	3
197424	1.0	a914ecef9c12fdb9bede64bb703d877	fast	4.0	6	3010	4
197425	1.0	a914ecef9c12fdb9bede64bb703d877	fast	4.0	5	1836	3
197426	1.0	c81e155d85dae5430a8cee6f2242e82c	sandwich	1.0	1	1175	1
197427	1.0	c81e155d85dae5430a8cee6f2242e82c	sandwich	1.0	4	2605	4

197428 rows × 15 columns

```
In [34]: df.isna().sum()
```

```
Out[34]: market_id      987
store_id      0
store_primary_category  4760
order_protocol    995
total_items      0
subtotal      0
num_distinct_items  0
min_item_price    0
max_item_price    0
total_onshift_partners  989
total_busy_partners    989
total_outstanding_orders  989
time_taken_minutes    7
order_hour      0
order_day_of_week    0
dtype: int64
```

```
In [35]: df[df["total_onshift_partners"].isnull()].dropna(inplace=True)
```

```
In [36]: df.isna().sum()
```

```
Out[36]: market_id      987
store_id      0
store_primary_category  4760
order_protocol    995
total_items      0
subtotal      0
num_distinct_items  0
min_item_price    0
max_item_price    0
total_onshift_partners  989
total_busy_partners    989
total_outstanding_orders  989
time_taken_minutes    7
order_hour      0
order_day_of_week    0
dtype: int64
```

```
In [37]: df = df[~df['total_onshift_partners'].isnull()]
```





```
In [38]: df.isna().sum()
```

```
Out[38]: market_id          0
store_id          0
store_primary_category  4268
order_protocol     508
total_items        0
subtotal           0
num_distinct_items  0
min_item_price     0
max_item_price     0
total_onshift_partners  0
total_busy_partners  0
total_outstanding_orders  0
time_taken_minutes  7
order_hour         0
order_day_of_week  0
dtype: int64
```

```
In [39]: df = df[~df['order_protocol'].isnull()]
```

```
In [40]: df.isna().sum()
```

```
Out[40]: market_id          0
store_id          0
store_primary_category  4005
order_protocol        0
total_items          0
subtotal             0
num_distinct_items    0
min_item_price        0
max_item_price        0
total_onshift_partners  0
total_busy_partners    0
total_outstanding_orders  0
time_taken_minutes     7
order_hour            0
order_day_of_week      0
dtype: int64
```

```
In [41]: df = df[~df['time_taken_minutes'].isnull()]
```

```
In [42]: df.isna().sum()
```

```
Out[42]: market_id          0
store_id          0
store_primary_category  4005
order_protocol        0
total_items          0
subtotal             0
num_distinct_items    0
min_item_price        0
max_item_price        0
total_onshift_partners  0
total_busy_partners    0
total_outstanding_orders  0
time_taken_minutes     0
order_hour            0
order_day_of_week      0
dtype: int64
```



```
In [43]: df[df["store_primary_category"].isna()]
```

Out[43]:

	market_id	store_id	store_primary_category	order_protocol	total_items	subtotal	num_distinct_items	
	2	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	1	1900	1
	3	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	6	6900	5
	4	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	3	3900	3
	5	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	3	5000	3
	6	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	2	3900	2
	...	...	...	...	...	...	...	...
	197208	1.0	77c493ec14246d748db3ee8fce0092db	NaN	1.0	7	5100	6
	197209	1.0	77c493ec14246d748db3ee8fce0092db	NaN	1.0	7	7200	6
	197210	1.0	77c493ec14246d748db3ee8fce0092db	NaN	1.0	3	2800	3
	197211	1.0	77c493ec14246d748db3ee8fce0092db	NaN	1.0	2	1400	2
	197212	1.0	77c493ec14246d748db3ee8fce0092db	NaN	1.0	5	2800	5

4005 rows × 15 columns

```
In [44]: df[df["store_id"]=="f0ade77b43923b38237db569b016ba25"]
```

Out[44]:

	market_id	store_id	store_primary_category	order_protocol	total_items	subtotal	num_distinct_items	min_
1	2.0	f0ade77b43923b38237db569b016ba25	mexican	2.0	1	1900	1	
2	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	1	1900	1	
3	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	6	6900	5	
4	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	3	3900	3	
5	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	3	5000	3	
6	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	2	3900	2	
7	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	4	4850	4	
8	2.0	f0ade77b43923b38237db569b016ba25	indian	3.0	4	4771	3	
9	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	2	2100	2	
10	3.0	f0ade77b43923b38237db569b016ba25	NaN	4.0	4	4300	4	
11	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	2	2200	2	
12	3.0	f0ade77b43923b38237db569b016ba25	NaN	1.0	1	1900	1	
13	3.0	f0ade77b43923b38237db569b016ba25	NaN	4.0	4	4986	4	

```
In [45]: df[df["store_primary_category"].isna()][["store_id"]].nunique()
```

Out[45]: 632

```
In [46]: df2=df[df["store_primary_category"].isna()][["store_id"]].unique()
```

# Imputing store\_primary\_category by mode



```
In [48]: # Function to impute missing values by mode, handling ties randomly
def impute_by_mode(df, column):
    # Get the mode(s)
    modes = df[column].mode()
    if len(modes) > 1:
        # If there are ties, choose one randomly with equal probability
        chosen_mode = np.random.choice(modes)
    else:
        # If no tie, use the single mode
        chosen_mode = modes[0]
    # Impute missing values with the chosen mode
    df[column].fillna(chosen_mode, inplace=True)
# List of columns to impute
columns_to_impute = ['store_primary_category']
# Apply the function to each column
for column in columns_to_impute:
    impute_by_mode(df, column)
```

```
In [49]: df.isna().sum()
```

```
Out[49]: market_id      0
store_id      0
store_primary_category  0
order_protocol  0
total_items    0
subtotal      0
num_distinct_items  0
min_item_price  0
max_item_price  0
total_onshift_partners  0
total_busy_partners  0
total_outstanding_orders  0
time_taken_minutes  0
order_hour      0
order_day_of_week  0
dtype: int64
```

```
In [50]: df.shape
```

```
Out[50]: (195924, 15)
```

```
In [51]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 195924 entries, 0 to 197427
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   market_id             195924 non-null float64
 1   store_id              195924 non-null object
 2   store_primary_category 195924 non-null object
 3   order_protocol        195924 non-null float64
 4   total_items           195924 non-null int64
 5   subtotal              195924 non-null int64
 6   num_distinct_items    195924 non-null int64
 7   min_item_price        195924 non-null int64
 8   max_item_price        195924 non-null int64
 9   total_onshift_partners 195924 non-null float64
10   total_busy_partners    195924 non-null float64
11   total_outstanding_orders 195924 non-null float64
12   time_taken_minutes    195924 non-null float64
13   order_hour            195924 non-null int64
14   order_day_of_week     195924 non-null int64
dtypes: float64(6), int64(7), object(2)
memory usage: 23.9+ MB
```

```
In [52]: store_name_counts = df['store_id'].value_counts()
df['store_name_enc'] = df['store_id'].map(store_name_counts)
```

```
In [54]: df = df.drop('store_name_enc', axis=1)
```




In [55]: df

Out[55]:

	market_id	store_id	store_primary_category	order_protocol	total_items	subtotal	num_distinct_items
0	1.0	df263d996281d984952c07998dc54358	american	1.0	4	3441	4
1	2.0	f0ade77b43923b38237db569b016ba25	mexican	2.0	1	1900	1
2	3.0	f0ade77b43923b38237db569b016ba25	american	1.0	1	1900	1
3	3.0	f0ade77b43923b38237db569b016ba25	american	1.0	6	6900	5
4	3.0	f0ade77b43923b38237db569b016ba25	american	1.0	3	3900	3
...	...	...	...	...	...	...	...
197423	1.0	a914ecef9c12fdb9bede64bb703d877	fast	4.0	3	1389	3
197424	1.0	a914ecef9c12fdb9bede64bb703d877	fast	4.0	6	3010	4
197425	1.0	a914ecef9c12fdb9bede64bb703d877	fast	4.0	5	1836	3
197426	1.0	c81e155d85dae5430a8cee6f2242e82c	sandwich	1.0	1	1175	1
197427	1.0	c81e155d85dae5430a8cee6f2242e82c	sandwich	1.0	4	2605	4

195924 rows × 15 columns



## Using Label Encoding for store name

In [56]: from sklearn.preprocessing import LabelEncoder


```
In [57]: label_encoder = LabelEncoder()
df['store_name_encoded'] = label_encoder.fit_transform(df['store_id'])
```

In [58]: df

Out[58]:

	market_id	store_id	store_primary_category	order_protocol	total_items	subtotal	num_distinct_items
0	1.0	df263d996281d984952c07998dc54358	american	1.0	4	3441	4
1	2.0	f0ade77b43923b38237db569b016ba25	mexican	2.0	1	1900	1
2	3.0	f0ade77b43923b38237db569b016ba25	american	1.0	1	1900	1
3	3.0	f0ade77b43923b38237db569b016ba25	american	1.0	6	6900	5
4	3.0	f0ade77b43923b38237db569b016ba25	american	1.0	3	3900	3
...	...	...	...	...	...	...	...
197423	1.0	a914ecef9c12fdb9bede64bb703d877	fast	4.0	3	1389	3
197424	1.0	a914ecef9c12fdb9bede64bb703d877	fast	4.0	6	3010	4
197425	1.0	a914ecef9c12fdb9bede64bb703d877	fast	4.0	5	1836	3
197426	1.0	c81e155d85dae5430a8cee6f2242e82c	sandwich	1.0	1	1175	1
197427	1.0	c81e155d85dae5430a8cee6f2242e82c	sandwich	1.0	4	2605	4

195924 rows × 16 columns



In [59]: df=df.drop("store\_id",axis=1)



In [61]: df

Out[61]:

	market_id	store_primary_category	order_protocol	total_items	subtotal	num_distinct_items	min_item_price	max_item_price	total
0	1.0	american	1.0	4	3441	4	557	1239	
1	2.0	mexican	2.0	1	1900	1	1400	1400	
2	3.0	american	1.0	1	1900	1	1900	1900	
3	3.0	american	1.0	6	6900	5	600	1800	
4	3.0	american	1.0	3	3900	3	1100	1600	
...	...	...	...	...	...	...	...	...	
197423	1.0	fast	4.0	3	1389	3	345	649	
197424	1.0	fast	4.0	6	3010	4	405	825	
197425	1.0	fast	4.0	5	1836	3	300	399	
197426	1.0	sandwich	1.0	1	1175	1	535	535	
197427	1.0	sandwich	1.0	4	2605	4	425	750	

195924 rows × 15 columns

```
In [62]: duplicates = df.duplicated()
# Print the original DataFrame with a marker for duplicates
print(df.loc[duplicates])
```

	market_id	store_primary_category	order_protocol	total_items	\
139263	6.0	indian	3.0	2	
166281	6.0	cafe	4.0	1	
	subtotal	num_distinct_items	min_item_price	max_item_price	\
139263	1650	1	825	825	
166281	350	1	350	350	
	total_onshift_partners	total_busy_partners	total_outstanding_orders	\	
139263	39.813559	40.40678	51.135593		
166281	39.813559	40.40678	51.135593		
	time_taken_minutes	order_hour	order_day_of_week	store_name_encoded	
139263	24.0	4	1	2637	
166281	39.0	4	4	1501	

In [63]: df=df.drop\_duplicates()

In [64]: df

Out[64]:

	market_id	store_primary_category	order_protocol	total_items	subtotal	num_distinct_items	min_item_price	max_item_price	total
0	1.0	american	1.0	4	3441	4	557	1239	
1	2.0	mexican	2.0	1	1900	1	1400	1400	
2	3.0	american	1.0	1	1900	1	1900	1900	
3	3.0	american	1.0	6	6900	5	600	1800	
4	3.0	american	1.0	3	3900	3	1100	1600	
...	...	...	...	...	...	...	...	...	
197423	1.0	fast	4.0	3	1389	3	345	649	
197424	1.0	fast	4.0	6	3010	4	405	825	
197425	1.0	fast	4.0	5	1836	3	300	399	
197426	1.0	sandwich	1.0	1	1175	1	535	535	
197427	1.0	sandwich	1.0	4	2605	4	425	750	

195922 rows × 15 columns



In [65]: `df.isna().sum()`

```
Out[65]: market_id          0
store_primary_category    0
order_protocol            0
total_items              0
subtotal                 0
num_distinct_items       0
min_item_price           0
max_item_price           0
total_onshift_partners    0
total_busy_partners       0
total_outstanding_orders  0
time_taken_minutes       0
order_hour               0
order_day_of_week        0
store_name_encoded        0
dtype: int64
```

In [66]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 195922 entries, 0 to 197427
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            195922 non-null float64
1   store_primary_category               195922 non-null object
2   order_protocol                      195922 non-null float64
3   total_items                         195922 non-null int64
4   subtotal                            195922 non-null int64
5   num_distinct_items                  195922 non-null int64
6   min_item_price                      195922 non-null int64
7   max_item_price                      195922 non-null int64
8   total_onshift_partners              195922 non-null float64
9   total_busy_partners                 195922 non-null float64
10  total_outstanding_orders            195922 non-null float64
11  time_taken_minutes                  195922 non-null float64
12  order_hour                         195922 non-null int64
13  order_day_of_week                   195922 non-null int64
14  store_name_encoded                  195922 non-null int32
dtypes: float64(6), int32(1), int64(7), object(1)
memory usage: 23.2+ MB
```

label Encoding store\_primary\_category

In [69]: `label_encoder = LabelEncoder()`  
`df['store_primary_category_enc'] = label_encoder.fit_transform(df['store_primary_category'])`

In [70]: `df=df.drop("store_primary_category",axis=1)`

In [71]: `df`

```
Out[71]:
```

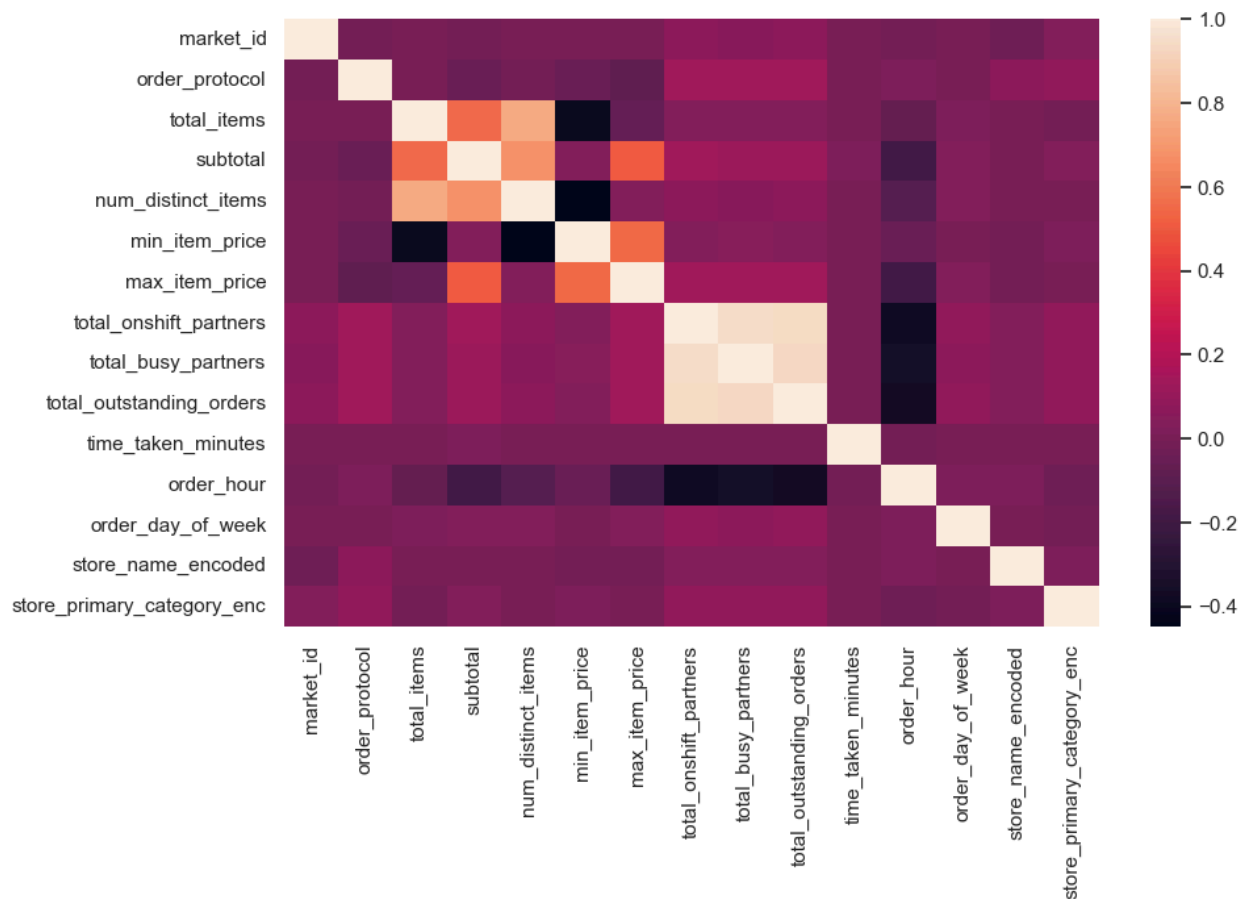
	market_id	order_protocol	total_items	subtotal	num_distinct_items	min_item_price	max_item_price	total_onshift_partners	total_busy_partners
0	1.0	1.0	4	3441	4	557	1239	33.0	
1	2.0	2.0	1	1900	1	1400	1400	1.0	
2	3.0	1.0	1	1900	1	1900	1900	1.0	
3	3.0	1.0	6	6900	5	600	1800	1.0	
4	3.0	1.0	3	3900	3	1100	1600	6.0	
...	...	...	...	...	...	...	...	...	
197423	1.0	4.0	3	1389	3	345	649	17.0	
197424	1.0	4.0	6	3010	4	405	825	12.0	
197425	1.0	4.0	5	1836	3	300	399	39.0	
197426	1.0	1.0	1	1175	1	535	535	7.0	
197427	1.0	1.0	4	2605	4	425	750	20.0	

195922 rows × 10 columns

## Data Visualisation



```
In [74]: sns.heatmap(df.corr())
plt.show()
```




```
In [75]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 195922 entries, 0 to 197427
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            195922 non-null float64
1   order_protocol                       195922 non-null float64
2   total_items                          195922 non-null int64
3   subtotal                             195922 non-null int64
4   num_distinct_items                  195922 non-null int64
5   min_item_price                      195922 non-null int64
6   max_item_price                      195922 non-null int64
7   total_onshift_partners              195922 non-null float64
8   total_busy_partners                 195922 non-null float64
9   total_outstanding_orders            195922 non-null float64
10  time_taken_minutes                  195922 non-null float64
11  order_hour                          195922 non-null int64
12  order_day_of_week                   195922 non-null int64
13  store_name_encoded                   195922 non-null int32
14  store_primary_category_enc          195922 non-null int32
dtypes: float64(6), int32(2), int64(7)
memory usage: 22.4 MB
```



```
In [77]: # Create the scatter plot
sns.scatterplot(x='time_taken_minutes', y='subtotal', data=df)
# Set the x-axis Limit
plt.xlim(0, 800)
plt.show()
```



```
In [80]: sns.scatterplot(x='total_items', y='subtotal', hue='num_distinct_items', palette='coolwarm', data=df)
plt.show()
```



```
In [81]: df3=df.copy()
```

```
In [82]: df3.shape
```

```
Out[82]: (195922, 15)
```





```
In [83]: df3=df3.drop("store_name_encoded",axis=1)
```

## Removing outliers using LOF


```
In [84]: from sklearn.neighbors import LocalOutlierFactor
import matplotlib.pyplot as plt
model1 = LocalOutlierFactor(contamination=0.05)
df3['lof_anomaly_score'] = model1.fit_predict(df3)
```

```
In [85]: print("number of outliers : ",(len(df3.loc[(df3['lof_anomaly_score'] == -1)])))
df3=df3.loc[(df3['lof_anomaly_score'] == 1)]
```

number of outliers : 9797

```
In [86]: df3.drop(['lof_anomaly_score'],axis=1,inplace=True)
```

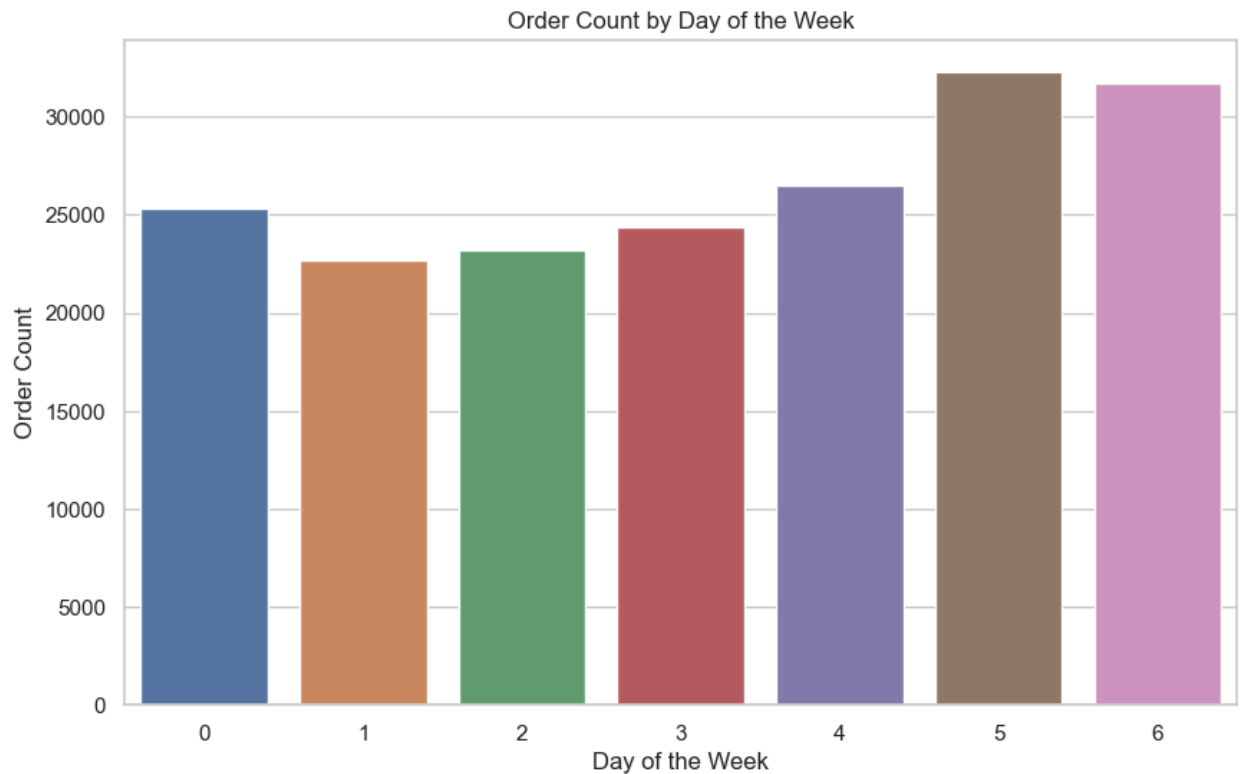
```
In [88]: # Create the scatter plot
sns.scatterplot(x='time_taken_minutes', y='subtotal', data=df3)
plt.show()
```




## Making various plots from features




```
In [89]: # Create a countplot for the 'order_day_of_week' column
sns.countplot(x='order_day_of_week', data=df3)
# Set the title and labels
plt.title('Order Count by Day of the Week')
plt.xlabel('Day of the Week')
plt.ylabel('Order Count')
# Show the plot
plt.show()
```




```
In [90]: # Create a countplot for the 'order_day_of_week' column
sns.countplot(x='order_hour', data=df3)
# Set the title and labels
plt.title('Order Count by Day of the Week')
plt.xlabel('Day of the Week')
plt.ylabel('Order Count')
# Show the plot
plt.show()
```




```
In [91]: sns.countplot(x=df.market_id)
plt.show()
```




```
In [92]: # Create a scatter plot for 'order_hour' vs 'time_taken_minutes'
plt.figure(figsize=(12, 6))
sns.scatterplot(x='order_hour', y='time_taken_minutes', data=df3)
# Set the title and labels
plt.title('Time Taken Minutes vs Order Hour')
plt.xlabel('Order Hour')
plt.ylabel('Time Taken Minutes')
# Show the plot
plt.show()
```




```
In [93]: plt.figure(figsize=(12, 6))  
sns.lineplot(x='total_onshift_partners', y='time_taken_minutes', data=df3, ci=None)  
plt.show()
```




```
In [95]: plt.figure(figsize=(12, 6))  
sns.lineplot(x='total_outstanding_orders', y='time_taken_minutes', data=df3, ci=None)  
plt.show()
```




```
In [96]: plt.figure(figsize=(12, 6))
sns.lineplot(x='total_busy_partners', y='time_taken_minutes', data=df3, ci=None)
plt.show()
```



```
In [97]: plt.figure(figsize=(12, 6))
sns.lineplot(x='order_hour', y='time_taken_minutes', data=df3, ci=None)
# Set the title and labels
plt.title('Average Time Taken Minutes vs Order Hour')
plt.xlabel('Order Hour')
plt.ylabel('Time Taken Minutes')
# Show the plot
plt.show()
```



```
In [98]: # Create a box plot for 'order_hour' vs 'time_taken_minutes'
plt.figure(figsize=(12, 6))
sns.boxplot(x='order_hour', y='time_taken_minutes', data=df3)
# Set the title and labels
plt.title('Distribution of Time Taken Minutes vs Order Hour')
plt.xlabel('Order Hour')
plt.ylabel('Time Taken Minutes')
# Show the plot
plt.show()
```



```
In [99]: y = df3['time_taken_minutes']
x = df3.drop(['time_taken_minutes'], axis=1)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
In [100]: x
```

```
Out[100]:
```

	market_id	order_protocol	total_items	subtotal	num_distinct_items	min_item_price	max_item_price	total_onshift_partners	total_i
0	1.0	1.0	4	3441	4	557	1239	33.0	
1	2.0	2.0	1	1900	1	1400	1400	1.0	
2	3.0	1.0	1	1900	1	1900	1900	1.0	
3	3.0	1.0	6	6900	5	600	1800	1.0	
4	3.0	1.0	3	3900	3	1100	1600	6.0	
...	...	...	...	...	...	...	...	...	...
197423	1.0	4.0	3	1389	3	345	649	17.0	
197424	1.0	4.0	6	3010	4	405	825	12.0	
197425	1.0	4.0	5	1836	3	300	399	39.0	
197426	1.0	1.0	1	1175	1	535	535	7.0	
197427	1.0	1.0	4	2605	4	425	750	20.0	

186125 rows × 13 columns



In [101]:

y

Out[101]:

```

0         62.0
1         67.0
2         29.0
3         51.0
4         39.0
...
197423    65.0
197424    56.0
197425    50.0
197426    65.0
197427    37.0
Name: time_taken_minutes, Length: 186125, dtype: float64

```

In [102]:

```

#random forest model training
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.metrics import mean_absolute_error
from sklearn.ensemble import RandomForestRegressor

```

## Creating baseline model RF to compare with Neural Networks

In [103]:

```

regressor = RandomForestRegressor()
regressor.fit(X_train, y_train)

```

Out[103]:

```

RandomForestRegressor
RandomForestRegressor()

```

In [104]:

```

prediction = regressor.predict(X_test)
mse = mean_squared_error(y_test, prediction)
rmse = mse**.5
print("mse : ", mse)
print("rmse : ", rmse)
mae = mean_absolute_error(y_test, prediction)
print('mae:', mae)

```

```

mse : 203.29293151608346
rmse : 14.25808302388801
mae: 10.898218183291227

```

In [105]:

```

r2_score(y_test, prediction)

```

Out[105]:

```

0.2670855268818707

```

In [107]:

```

def MAPE(Y_actual, Y_Predicted):
    mape = np.mean(np.abs((Y_actual - Y_Predicted)/Y_actual))*100
    return mape

```

In [108]:

```

print("mape : ", MAPE(y_test, prediction))

```

```


mape : 26.222978362229632

```



```
In [109]: sorted_idx = regressor.feature_importances_.argsort()
plt.barh(df3.columns[sorted_idx], regressor.feature_importances_[sorted_idx])
plt.xlabel("Random Forest Feature Importance")
```

```
Out[109]: Text(0.5, 0, 'Random Forest Feature Importance')
```



## Train-Test Splitting Standard Scaling

```
In [110]: from sklearn import preprocessing
from sklearn.model_selection import train_test_split

# Initialize the MinMaxScaler
scaler = preprocessing.MinMaxScaler()

# Fit and transform the data
x_scaled = scaler.fit_transform(x)

# Split the scaled data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.2, random_state=42)
```

## Creating Neural Network Architecture

```
In [111]: model = Sequential()
model.add(Dense(11, kernel_initializer='normal'))
model.add(Dense(256, activation='relu'))
model.add(Dense(512, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(1, activation='linear'))
```

## Model Training





```
In [112]: from tensorflow.keras.optimizers import Adam

adam = Adam(learning_rate=0.01)
model.compile(loss='mse', optimizer=adam, metrics=['mse', 'mae'])
history = model.fit(X_train, y_train, epochs=30, batch_size=512, verbose=1, validation_data=(X_test, y_test))
```




```
Epoch 1/30
291/291 [=====] - 6s 10ms/step - loss: 282.6307 - mse: 282.6307 - mae: 12.6141 - val_loss: 231.1850 - val_mse: 231.1850 - val_mae: 11.1970
Epoch 2/30
291/291 [=====] - 3s 9ms/step - loss: 224.8160 - mse: 224.8160 - mae: 11.4173 - val_loss: 217.4656 - val_mse: 217.4656 - val_mae: 11.2839
Epoch 3/30
291/291 [=====] - 3s 10ms/step - loss: 224.3059 - mse: 224.3059 - mae: 11.3947 - val_loss: 220.9363 - val_mse: 220.9363 - val_mae: 11.6191
Epoch 4/30
291/291 [=====] - 3s 10ms/step - loss: 224.5054 - mse: 224.5054 - mae: 11.3996 - val_loss: 220.7947 - val_mse: 220.7947 - val_mae: 11.6464
Epoch 5/30
291/291 [=====] - 3s 10ms/step - loss: 221.8747 - mse: 221.8747 - mae: 11.3358 - val_loss: 214.4131 - val_mse: 214.4131 - val_mae: 11.0570
Epoch 6/30
291/291 [=====] - 3s 10ms/step - loss: 220.3652 - mse: 220.3652 - mae: 11.2856 - val_loss: 216.2168 - val_mse: 216.2168 - val_mae: 11.0922
Epoch 7/30
291/291 [=====] - 3s 9ms/step - loss: 220.2037 - mse: 220.2037 - mae: 11.2854 - val_loss: 213.7458 - val_mse: 213.7458 - val_mae: 11.2511
Epoch 8/30
291/291 [=====] - 3s 9ms/step - loss: 221.1623 - mse: 221.1623 - mae: 11.3061 - val_loss: 213.0859 - val_mse: 213.0859 - val_mae: 11.1524
Epoch 9/30
291/291 [=====] - 3s 10ms/step - loss: 219.3359 - mse: 219.3359 - mae: 11.2506 - val_loss: 223.1522 - val_mse: 223.1522 - val_mae: 11.7961
Epoch 10/30
291/291 [=====] - 3s 10ms/step - loss: 218.5172 - mse: 218.5172 - mae: 11.2343 - val_loss: 213.8713 - val_mse: 213.8713 - val_mae: 11.3567
Epoch 11/30
291/291 [=====] - 3s 10ms/step - loss: 218.3957 - mse: 218.3957 - mae: 11.2331 - val_loss: 238.1378 - val_mse: 238.1378 - val_mae: 12.4312
Epoch 12/30
291/291 [=====] - 3s 11ms/step - loss: 218.2340 - mse: 218.2340 - mae: 11.2343 - val_loss: 220.8149 - val_mse: 220.8149 - val_mae: 11.7375
Epoch 13/30
291/291 [=====] - 3s 9ms/step - loss: 217.8660 - mse: 217.8660 - mae: 11.2131 - val_loss: 211.0166 - val_mse: 211.0166 - val_mae: 10.9998
Epoch 14/30
291/291 [=====] - 3s 9ms/step - loss: 217.0804 - mse: 217.0804 - mae: 11.1905 - val_loss: 209.9097 - val_mse: 209.9097 - val_mae: 11.1757
Epoch 15/30
291/291 [=====] - 3s 10ms/step - loss: 213.8961 - mse: 213.8961 - mae: 11.0969 - val_loss: 207.6500 - val_mse: 207.6500 - val_mae: 10.9583
Epoch 16/30
291/291 [=====] - 3s 10ms/step - loss: 213.3959 - mse: 213.3959 - mae: 11.0918 - val_loss: 222.3456 - val_mse: 222.3456 - val_mae: 11.8455
Epoch 17/30
291/291 [=====] - 3s 11ms/step - loss: 214.4584 - mse: 214.4584 - mae: 11.1235 - val_loss: 206.4343 - val_mse: 206.4343 - val_mae: 10.8411
Epoch 18/30
291/291 [=====] - 3s 11ms/step - loss: 212.3076 - mse: 212.3076 - mae: 11.0493 - val_loss: 206.4986 - val_mse: 206.4986 - val_mae: 10.8529
Epoch 19/30
291/291 [=====] - 3s 9ms/step - loss: 213.4192 - mse: 213.4192 - mae: 11.0920 - val_loss: 205.9552 - val_mse: 205.9552 - val_mae: 10.7876
Epoch 20/30
291/291 [=====] - 3s 10ms/step - loss: 211.1879 - mse: 211.1879 - mae: 11.0295 - val_loss: 207.3801 - val_mse: 207.3801 - val_mae: 11.1824
Epoch 21/30
291/291 [=====] - 3s 10ms/step - loss: 210.0776 - mse: 210.0776 - mae: 11.0004 - val_loss: 207.7153 - val_mse: 207.7153 - val_mae: 11.1545
Epoch 22/30
291/291 [=====] - 3s 10ms/step - loss: 210.1328 - mse: 210.1328 - mae: 10.9997 - val_loss: 204.5477 - val_mse: 204.5477 - val_mae: 10.9081
Epoch 23/30
291/291 [=====] - 3s 10ms/step - loss: 211.6617 - mse: 211.6617 - mae: 11.0298 - val_loss: 207.4345 - val_mse: 207.4345 - val_mae: 10.8461
Epoch 24/30
291/291 [=====] - 3s 10ms/step - loss: 211.8860 - mse: 211.8860 - mae: 11.0475 - val_loss: 205.8262 - val_mse: 205.8262 - val_mae: 11.0192
Epoch 25/30
291/291 [=====] - 3s 10ms/step - loss: 210.1552 - mse: 210.1552 - mae: 10.9974 - val_loss: 206.0043 - val_mse: 206.0043 - val_mae: 10.8366
Epoch 26/30
291/291 [=====] - 3s 10ms/step - loss: 209.6816 - mse: 209.6816 - mae: 10.9877 - val_loss: 209.1212 - val_mse: 209.1212 - val_mae: 10.7518
Epoch 27/30
291/291 [=====] - 3s 10ms/step - loss: 208.9299 - mse: 208.9299 - mae: 10.9707 - val_loss: 204.7607 - val_mse: 204.7607 - val_mae: 10.7839
Epoch 28/30
291/291 [=====] - 3s 10ms/step - loss: 209.0205 - mse: 209.0205 - mae: 10.9707 - val_loss: 205.2039 - val_mse: 205.2039 - val_mae: 11.0781
```



Epoch 29/30  
 291/291 [=====] - 3s 10ms/step - loss: 208.6797 - mse: 208.6797 - mae: 10.9603 - val\_loss: 203.5120 - val\_mse: 203.5120 - val\_mae: 10.8419  
 Epoch 30/30  
 291/291 [=====] - 3s 9ms/step - loss: 208.4299 - mse: 208.4299 - mae: 10.9457 - val\_loss: 203.7377 - val\_mse: 203.7377 - val\_mae: 10.8553

## Comparing losses with epochs

```
In [113]: def plot_history(history, key):
plt.plot(history.history[key])
plt.plot(history.history['val_'+key])
plt.xlabel("Epochs")
plt.ylabel(key)
plt.legend([key, 'val_'+key])
plt.show()
# Plot the history
plot_history(history, 'loss')
```



```
In [114]: z = model.predict(X_test)
```

1164/1164 [=====] - 2s 2ms/step

```
In [115]: r2_score(y_test, z)
```

Out[115]: 0.26548217263663154

## MAE RMSE MSE values for Neural Networks

```
In [116]: mse = mean_squared_error(y_test, z)
rmse = mse**.5
print("mse : ",mse)
print("rmse : ",rmse)
mae = mean_absolute_error(y_test, z)
print("mae : ",mae)
```

mse : 203.7376636051998  
 rmse : 14.27367029201669  
 mae : 10.855294681751463

In [ ]:



